

## Overall Goal for the Month:

To develop a framework and architectural design for the Context-Aware Spatial Description (CAS-D) system and implement a functional MVP codebase that can process a video, manage contextual memory, and generate a description.

---

## Week 1: Foundations & Environment Setup

Objective: Establish the theoretical basis and create the complete data processing pipeline and project structure, which is often the most time-consuming part of implementation.

- Theoretical Work:
  - a. In-depth Paper Analysis: Deconstruct the core mechanisms of MC-ViT (non-parametric memory) and Video-3D LLM (position-aware representations). Diagram their data flows.
  - b. Problem Formalization: Write a clear problem statement. Why is combining these two ideas novel and necessary for rich, contextual 3D scene description?
  - c. Initial Architecture Sketch: Draw a rough, first-pass diagram of the CAS-D framework on a whiteboard or paper.
- Coding Work (MVP Focus: Data Pipeline):
  - a. Environment Setup:
    - i. Set up a Python virtual environment (e.g., Conda).
    - ii. Install essential libraries: torch, torchvision, transformers, timm (for ViT models), numpy, and a library for handling 3D data if needed (e.g., open3d).
  - b. Project Scaffolding: Create the basic file structure:
  - c. Generated code (tentative structure)

```
casd_project/
├── data/           # For dataset files
├── src/
│   ├── dataset.py  # Data loading and preprocessing
│   ├── model.py    # The main CAS-D model architecture
│   ├── agent.py    # The memory consolidation agent
│   ├── train.py    # Training and evaluation loop
│   └── config.py   # Hyperparameters and settings
├── notebooks/     # For exploration and visualization
├── README.md
└── requirements.txt
```

- d. Dataset & DataLoader:

- i. Download a suitable dataset. ScanNet is ideal because it provides RGB-D video streams with camera poses.
- ii. In `src/dataset.py`, implement a PyTorch Dataset class that can:
  - Load a video sequence from the dataset.
  - For each frame, retrieve the RGB image, depth map, and camera intrinsic/extrinsic parameters.
  - Implement the "Maximum Coverage Sampling" from Video-3D LLM as a preprocessing step to select a fixed number of

Deliverable for Week 1:

- Theory: A document with diagrams and summaries of the key papers and a finalized problem statement.
  - Code: A functional data pipeline. You should be able to run a script that loads a batch of sampled video frames and their corresponding 3D metadata, ready to be fed into a model. This is a huge milestone.
- 

## Week 2: Position-Aware Encoding

Objective: To implement the "eyes" of your system—the module that turns video frames into 3D-aware representations—and refine the system's architectural design.

- Theoretical Work:
  1. Formal Architecture Design: Turn your initial sketch into a formal architectural diagram with all modules, inputs, outputs, and tensor shapes clearly defined. This will be guided by your coding work.
  2. Mathematical Formalism: Write down the exact equations for creating the position-aware representation ( $e_{\text{vis}} = e_{\text{img}} + e_{\text{coord}}$ ), including the 3D back-projection and the sinusoidal position encoding.
- Coding Work (MVP Focus: The Encoder):
  1. Implement the Encoder Module: In `src/model.py`, create a `PositionAwareEncoder` module.
    - Instantiate a pre-trained Vision Transformer (e.g., ViT-Base from `timm` or Hugging Face).
    - Write a helper function that takes a batch of depth maps and camera parameters and calculates the 3D world coordinates for the center of each patch.
    - Implement the sinusoidal position encoding function for the  $(x, y, z)$  coordinates.
    - Modify the forward pass:
      - Get patch embeddings ( $e_{\text{img}}$ ) from the ViT.
      - Calculate coordinate embeddings ( $e_{\text{coord}}$ ).
      - Return their sum ( $e_{\text{vis}}$ ).
  2. Unit Testing: Write a simple test to ensure your encoder, given a dummy batch of data from your `DataLoader`, produces an output tensor of the correct shape (e.g., `[batch_size, num_patches, embedding_dim]`) without errors.

Deliverable for Week 2:

- Theory: A polished architectural diagram and the formalized mathematical definitions of your core components.
  - Code: A functional PositionAwareEncoder module that is tested and integrated with your data pipeline.
- 

## Week 3: The Memory Consolidation Agent

Objective: To implement the "brain" of your system—the agent that manages memory—and design a robust evaluation strategy to prove its effectiveness.

- Theoretical Work:
  1. Agent Strategy: Justify the choice of memory consolidation. For the MVP, k-means is an excellent choice. Theorize why k-means centroids of position-aware tokens would be a powerful form of memory (e.g., they represent persistent objects or spatial regions).
  2. Evaluation Design: Design a specific, novel evaluation task for your framework. Create 5-10 example Question/Answer pairs that are impossible to answer without long-term context (e.g., "What object was on the table before the person put the laptop there?").
- Coding Work (MVP Focus: The Agent & Memory):
  1. Implement the Agent: In `src/agent.py`, create a `MemoryAgent` class.
    - Implement the `consolidate` method. For the MVP, this method will take a set of token embeddings (from a past video segment) and use a simple k-means algorithm (e.g., a PyTorch implementation or a wrapper around scikit-learn's) to find K centroids.
  2. Integrate Agent into Model: In `src/model.py`:
    - Instantiate the `MemoryAgent`.
    - Add a persistent buffer to your `CASD_Model` to act as the `memory_bank`.
    - Modify the main forward pass to process video in segments. For each segment, it should use the agent to update the `memory_bank` based on the previous segment's encodings.
  3. Decoder with Cross-Attention:
    - Choose a pre-trained decoder-only LLM (e.g., GPT-2, Llama).
    - In the forward pass, feed the decoder two things:
      1. The embeddings for the current segment (`e_vis_t`).
      2. The `memory_bank` as the `encoder_hidden_states` argument. This enables cross-attention.

Deliverable for Week 3:

- Theory: A document justifying the memory strategy and detailing the novel, context-dependent evaluation protocol.
- Code: A model that can process a video segment, update a memory bank via the k-means agent, and is wired for cross-attention.

---

## Week 4: Integration, Training, and Final Proposal

Objective: To connect all components, run a proof-of-concept training loop, and synthesize all work into a final project proposal or paper draft.

- Theoretical Work:
  1. Write the Full Proposal/Paper Draft:
    - Write the Introduction and Related Work sections.
    - Write the Methodology section, using your diagrams and formalisms from previous weeks.
    - Write the Experiments section, detailing your dataset, the MVP implementation, and the proposed evaluation strategy.
- Coding Work (MVP Focus: End-to-End Run):
  1. Implement the Training Loop: In `src/train.py`:
    - Write a full training loop that iterates through your `DataLoader`.
    - In the loop, pass the data through your `CASD_Model`.
    - The model will output language logits from the decoder.
    - Compute the standard cross-entropy loss between the model's predictions and the ground-truth text descriptions.
    - Implement the backpropagation (`loss.backward()`) and optimizer step (`optimizer.step()`).
  2. Proof-of-Concept Run:
    - Run your `train.py` script on a very small subset of the data (e.g., 1-2 videos, batch size of 1) for a few dozen steps.
    - Goal: Verify that the code runs end-to-end without crashing and that the training loss decreases. This proves the entire architecture is viable.
  3. Code Finalization: Clean up the code, add comments, docstrings, and a comprehensive `README.md` explaining how to set up the environment and run the proof-of-concept.

Deliverable for Week 4:

- Theory: A polished, near-complete research proposal or paper draft.
- Code: A functional, end-to-end MVP codebase with a `README.md`. A successful demonstration shows the loss going down on a toy example.